

CS307

Database Principles

Stéphane Faroult
sfaroult@roughsea.com

朱悦铭 Zhu Yueming zhuym@sustc.edu.cn

UNPREDICTABLE is worse than SLOW

We have seen that because bind variables are checked when parsing and their values used to decide the execution plan, it can lead to plan instability when a query has aged out of the cache and is reparsed with different parameters.

This is why some DBA may try (when the DBMS allows it) to "attach" an execution plan known to be OK to a query. This is a short-term fix.

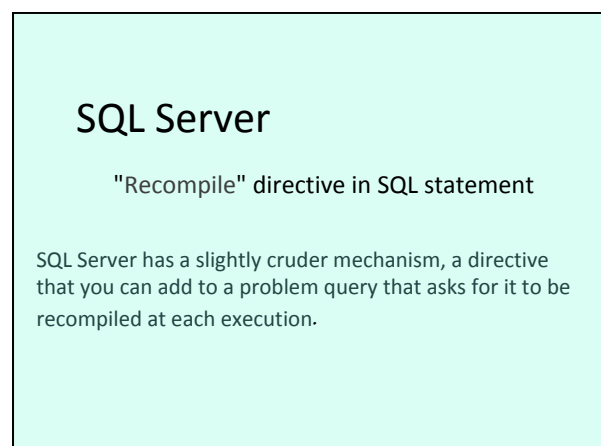
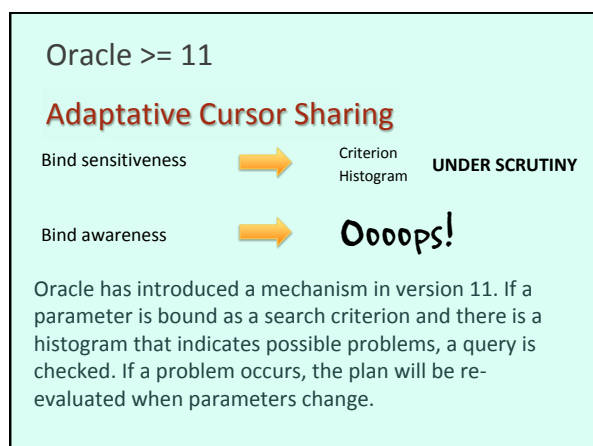
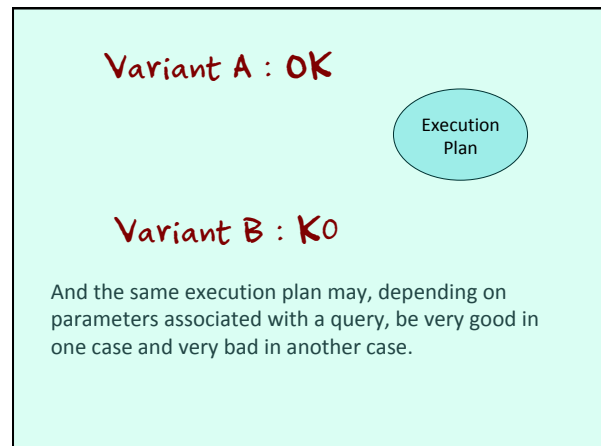
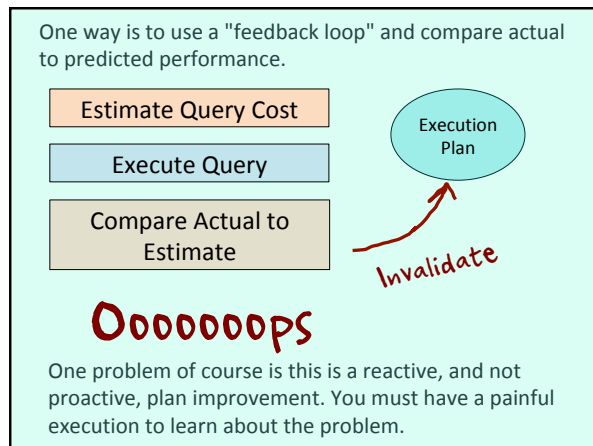
"Attach" execution plan to query



What can the optimizer do?



Let's see how DBMS vendors try to address the problem.



ORACLEDynamic sampling 0 ... **2** ... 10

Occurs during hard-parse

When set to 3 and above tries
to validate guesses

An interesting, and often effective, feature in Oracle is "dynamic sampling" which basically asks the optimizer to guess less and check data a bit more.

A Developer perspective

Full control of the application

Let's now see performance from the perspective of the developer. The developer CAN change the application. One problem is that, if DBAs are often seasoned professionals, in many cases developers are the cheapest beginners that could be found.

Because too many applications are written carelessly, a lot of time is devoted to optimizing.

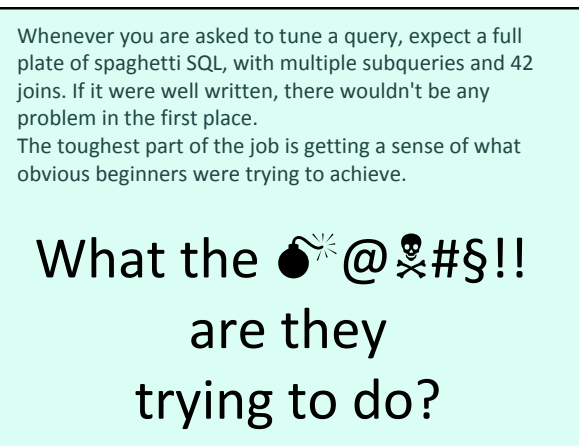

OPTIMIZING

Getting the best out of
a **given** situation.

There is only so much than you can do. If the design is bad, you may be able to improve things here and there, but not necessarily enough.

CORRECTLY WRITING (OR REWRITING) AN SQL QUERY

I have been asked umpteen time to help with SQL.



Remove what doesn't

shape

the result set

What is really important is what determines the magnitude of the number of rows returned by the query. When you rewrite a big, ugly query, first get rid of everything that changes nothing or little to the final number of rows returned.

Joins to tables without any

col = constant
condition

When you join to tables to which no direct criteria are applied, are these joins really important?

foreign key
or **NOT**
foreign key ?

If there is a foreign key, probably not: you know that for every row you'll find a match in the other table, so the join won't affect the number of rows returned.

```
select ...  
from TA, TB  
where ...  
    and TA.C = TB.C  
...
```

No other condition
on TB

If C is a **foreign key** for A that references B,
TB DOESN'T belong to the core of the query

If some values of C in TA **ARE NOT** in TB
TB BELONGS to the core of the query

In that case the join filters rows.

```

select distinct
  cons.id,
  coalesce(cons.definite_code, cons.provisional_code) dc,
  cons.name,
  cons.supplier_code,
  cons.registration_date,
  col.rco_name
from weighing w
  inner join production prod
    on prod.id = w.id
    inner join process_status prst
      on prst.prst_id = prod.prst_id
  left outer join composition comp
    on comp.formula_id = w.formula_id
  inner join constituent Cons
    on cons.id = w.id
  left outer join cons_color col
    on col.rcolor_id = cons.rcolor_id
where prod.usr_id = :userid
and prst.prst_code = 'PENDING'

```

Real life example. Two search criteria on two distinct tables, joined to each other..

```

select distinct
  cons.id,
  coalesce(cons.definite_code, cons.provisional_code) dc,
  cons.name,
  cons.supplier_code,
  cons.registration_date,
  col.rco_name
from weighing w
  inner join production prod
    on prod.id = w.id
    inner join process_status prst
      on prst.prst_id = prod.prst_id
  left outer join composition comp
    on comp.formula_id = w.formula_id
  inner join constituent cons
    on cons.id = w.id
  left outer join cons_color col
    on col.rcolor_id = cons.rcolor_id
where prod.usr_id = :userid
and prst.prst_code = 'PENDING'

```

Most returned values from one table, joined to the others through "w".

```

select distinct
  cons.id,
  coalesce(cons.definite_code, cons.provisional_code) dc,
  cons.name,
  cons.supplier_code,
  cons.registration_date,
  col.rco_name
from weighing w
  inner join production prod
    on prod.id = w.id
    inner join process_status prst
      on prst.prst_id = prod.prst_id
  left outer join composition comp
    on comp.formula_id = w.formula_id
  inner join constituent cons
    on cons.id = w.id
  left outer join cons_color col
    on col.rcolor_id = cons.rcolor_id
where prod.usr_id = :userid
and prst.prst_code = 'PENDING'

```

The only other column is returned through a left join.
Won't change the number of rows. Can go for now (we'll reinject it in the query after everything else is OK).

```

select distinct
  cons.id,
  coalesce(cons.definite_code,
    cons.provisional_code) dc,
  cons.name,
  cons.supplier_code,
  cons.registration_date from weighing w
  inner join production prod
    on prod.id = w.id
    inner join process_status prst
      on prst.prst_id = prod.prst_id
  left outer join composition comp
    on comp.formula_id = w.formula_id
  inner join constituent cons
    on cons.id = w.id
where prod.usr_id = :userid
and prst.prst_code = 'PENDING'

```

As a first step we'll consider this to be the core of the query.

Optimize the core of the query


Now let's try to retrieve all the rows as fast as possible.

Phase 2

```
select distinct
  cons.id,
  coalesce(cons.definite_code,
           cons.provisional_code) dc,
  cons.name,
  cons.supplier_code,
  cons.registration_date
from weighing w
  inner join production prod
    on prod.id = w.id
    inner join process_status prst
      on prst.prst_id = prod.prst_id
  left outer join composition comp
    on comp.formula_id = w.formula_id
  inner join constituent cons
    on cons.id = w.id
where prod.usr_id = :userid
   and prst.prst_code = 'PENDING'
```

This is how tables are related (dashed line means outer join)

Chain of tables



It can be represented as a "chain" of tables linked to each other through joins.

Tables from which data is returned
With or without conditions

Tables from which **NO** data is returned
Conditions only

Glue Tables
Join other tables

Classifying tables

```

select distinct
  cons.id,
  coalesce(cons.definite_code,
    cons.provisional_code) dc,
  cons.name,
  cons.supplier_code,
  cons.registration_date
from weighing w
  inner join production prod
    on prod.id = w.id
    inner join process_status prst
      on prst.prst_id = prod.prst_id
  left outer join composition comp
    on comp.formula_id = w.formula_id
  inner join constituent cons
    on cons.id = w.id
where prod.usr_id = :userid
and prst.prst_code = 'PENDING'

```

"comp" cannot be put into ANY category.
Additionally, DISTINCT smells of bad join.

Check outer (external) joins

NULL cannot be equal

NULL cannot be different

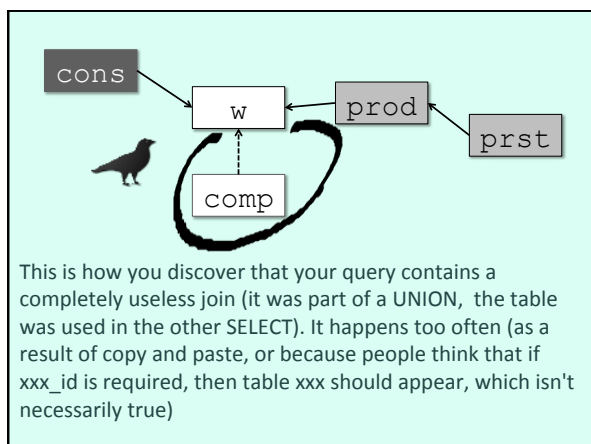
`outer.column = something`
without

or `outer.column is null`



but here we have no condition
whatsoever on "comp"!!

OUTER JOIN should be INNER JOIN



Tables from which data is returned
With or without conditions

Keep as is if it belongs to the core

Tables from which **NO** data is returned
Conditions only

Turn into a subquery

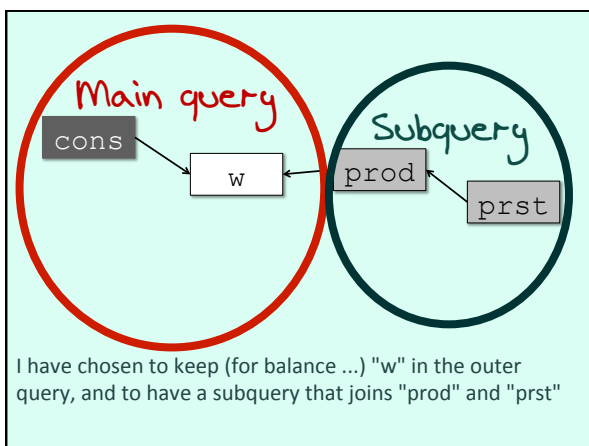
Subqueries (especially `IN ()` that takes care of duplicates) are a good way to get rid of `DISTINCT`, which should always be a black flag at the top of a big query. You may decide later which type of subquery you need and whether you should turn them into joins but they are a great way to make a query more legible.

Glue Tables
Join other tables

Useless at the end of a chain

Main query or subquery

A "glue" table that is glued to one useful table and nothing else can be safely disposed of, as we did with "comp". Remaining glue tables that link outer query tables to inner query tables can indifferently appear in the inner or outer query.



```


Select cons.id,
       coalesce(cons.definite_code,
                cons.prst_id) dc,
       cons.name,
       cons.supplier_code,
       cons.registration_date
from weighing w
   inner join constituent cons
           on cons.id = w.id
where w.id in
      (select prod.id
       from production prod
         inner join process_status prst
           on prst.prst_id = prod.prst_id
       where prod.usr_id = :userid
         and prst.prst_code = 'PENDING')

```

And lo, `DISTINCT` is gone. Is "w" really necessary? There was no FK, I couldn't tell, I kept it.

Two queries (**union**)

> 1 minute



0.4 seconds



The same kind of purely logical analysis applied to the other SELECT in the UNION (but the long one was this one), and speed improved by a factor of more than 100... No black magic.

Picture by MShades

Start with what

CUTS THROUGH ROWS



If you want your queries to run fast, you should first identify the criteria that set the order of magnitude of the number of rows in the result set.

Flickr: Leonid Mamchenkov

GROUP and sort
AS LITTLE AS POSSIBLE



Then use aggregates while only working with identifiers and "small" data, not big, fully joined rows.

Flickr: Belinda Hankins Miller

Join Late!



Keep for the very end (the top level, the outer query) all joins that change nothing (or very little) to the size of the result set.

Flickr: Jim Winstead

What about Hints?

The approach that is advocated here is based on the understanding of the functional aspect of a query, and focusing on a logical data search. Many self-styled experts take another approach, consider that the optimizer had it wrong (often true) and prefer telling the optimizer what to do using "hints".

Hint: special directive to tell the optimizer what to do

Those "hints" exist with all DBMS products and were mostly introduced when query optimizers were young; they were a convenient way to override optimizer bugs.



```
SELECT p.Name, pr.ProductReviewID
FROM Production.Product AS p
LEFT OUTER HASH JOIN
    Production.ProductReview AS pr
    ON p.ProductID = pr.ProductID
ORDER BY ProductReviewID DESC;
```

For instance you can tell to SQL Server which join algorithm it should use.



```
SELECT c.LastName, c.FirstName, e.Title
FROM HumanResources.Employee AS e
WITH (NOLOCK, INDEX (PK_Employee_EmployeeID))
JOIN Person.Contact AS c
    ON e.ContactID = c.ContactID
WHERE e.ManagerID = 3;
```

Or which index it should use.



```
SELECT * FROM table1
  USE INDEX (col1_index,col2_index)
WHERE col1=1
   AND col2=2
   AND col3=3;
```

You can do the same with MySQL.



With Oracle, you can even have more hints than SQL.


```
SELECT /*+ LEADING(e2 e1) USE_NL(e1)
          INDEX(e1 emp_emp_id_pk)
          USE_MERGE(j) FULL(j) */
       e1.first_name, e1.last_name,
       j.job_id, sum(e2.salary) total_sal
FROM employees e1, employees e2, job_history j
WHERE e1.employee_id = e2.manager_id
   AND e1.employee_id = j.employee_id
   AND e1.hire_date = j.start_date
GROUP BY e1.first_name, e1.last_name, j.job_id
ORDER BY total_sal;
```

/*+ Hints suck */

Hints are based on the assumption:

- 1) That you know better than the optimizer (in fact, it has far more information than you have about data and its storage)
- 2) That what is best now always will be.

Optimizer goal



We want to stay on the upper curve, everywhere.

I don't care about
execution plans
I care about correctness
and **SPEED**

Whether the DBMS is scanning, using indexes, nested loops, hash joins, is irrelevant as long as it's fast.


What the optimizer does is taking your query, testing different "starting points", possibly rewriting it in an equivalent way.



Then it finds something more efficient than plain "as is" execution, and runs it.

Of course sometimes the optimizer can be spectacularly off.

Two things can go wrong



Sometimes the optimizer simply makes a wrong choice.

99.99% of cases: Statistics issue

In most cases, it failed to make a correct cardinality estimate (how many rows were returned by a step), for reasons we have seen.
Feeding it better information will help.

The other case comes from the fact that the optimizer has little time to do its work (users are waiting for results). If your query is badly written, the optimizer may not be able even to contemplate something efficient.



Much better

Some people have come to consider that optimizers are lousy pieces of software that come in their way. They know better and should tell the optimizer what to do.



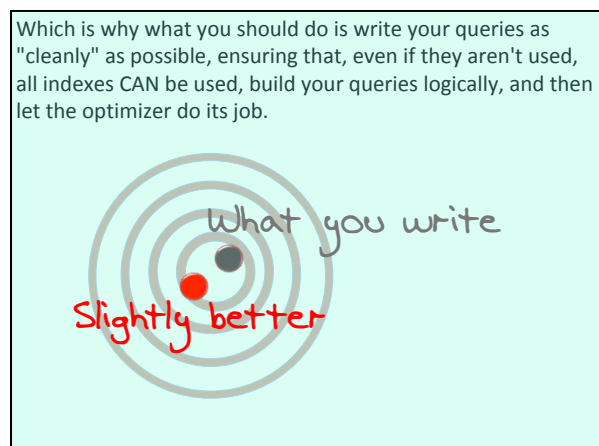



Picture by Dazhu Pagla

Hence "hints" that will direct the efforts of the optimizer and sometimes bring impressive improvement.

What you write

Much better

`/*+ hint */`





Jonathan Lewis's rules on Oracle hinting



1. Don't
2. If you must use hints, then assume you've used them incorrectly.
3. On every patch or upgrade to Oracle, assume every piece of hinted SQL is going to do the wrong thing ... because of (2) above. You've been lucky so far, but the patch/upgrade lets you discover your mistake.
4. Every time you apply some DDL to an object that appears in a piece of hinted SQL assume that the hinted SQL is going to do the wrong thing ... because of (2) above. You've been lucky so far, but the structural change lets you discover your mistake.

Jonathan is a highly respected British Oracle specialist.



This query says "for every (empid, emprcd) return the first row you find for which effdt is before today"

```
from ..., T_JOB J, ...
where ...
  and J.rowid = ( select /** INDEX_DESC(X, I_JOB)*/
                  X.rowid
                  from T_JOB X
                  where X.empid = J.empid
                  and X.emprcd = J.emprcd
                  and X.effdt <= SYSDATE
                  and rownum = 1 )
I_JOB on (empid, emprcd, effdt, effseq)
```

Diagram showing the hint **HIGHEST** applied to the **effdt** column, and another **HIGHEST** applied to the **effseq** column.

BECAUSE OF THE HINT, we return the highest value for effseq for the latest effdt, WHICH IS WHAT IS WANTED.

What happens if:

Index renamed

Index redefined

Optimizer changes



Flickr: Tony Crider

Technically speaking, a hint is a comment. A missing index won't make the query fail.

Hint silently ignored

The query will still do as instructed: return one row for which effdt is before today. But will it be the highest effdt before today and the highest effseq value for that effdt?

Wrong result

Or perhaps not. Who knows.

```
...
from ..., T_JOB J, ...
where ...
  and J.effdt = (select max(X.effdt)
                from T_JOB X
                where X.empid = J.empid
                   and X.emprcd = J.emprcd
                   and X.effdt <= SYSDATE)

  and J.effseq = (select max(Y.effseq)
                  from T_JOB Y
                  where Y.empid = J.empid
                     and Y.emprcd = J.emprcd
                     and Y.effdt = J.effdt)
```



This states exactly what we want; except that two correlated subqueries, one dependent on the result of the other, kill performances.

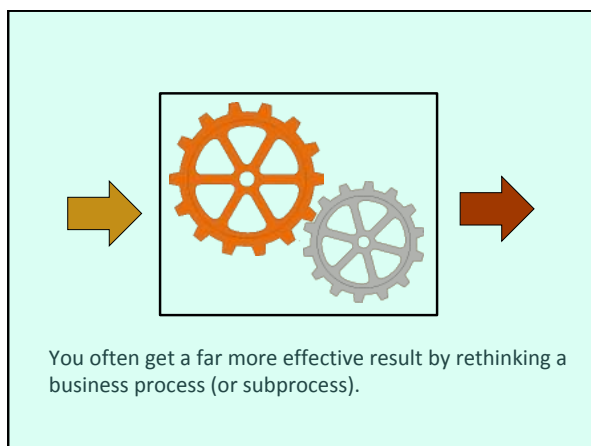

```
...
from ...,
  (select ...,
    rank() over (partition by empid,
                  emprcd
                  order by effdt desc,
                  effseq desc)
    rnk
  from T_JOB
  where effdt <= SYSDATE) J, ...
where ...
  and J.rnk = 1
```



But there are other, efficient ways to write it, which don't rely on side-effects to return the correct result.

GOING FARTHER THAN QUERIES

Most database specialists focus on trying to improve queries that seem to be bottlenecks. While it's often a good first step, you can go much further and I'd like to illustrate it.



Example

This example was made up, but based upon a real bank procedure, the purpose of which was to flag big money transactions as suspicious, in order to be able to investigate more and see whether they weren't related to cases of money laundering.

The process is based on thresholds, the amount of which depends on the currency used for the transactions. Suspicious transactions are logged and the amounts converted to a common currency.

```

ACCOUNTS
ACCOUNTID  NUMBER
AREAID     NUMBER

TRANSACTIONS
TXID       NUMBER      2,000,000 rows
TXDATE     DATE
MONEY_AMOUNT NUMBER
CURRENCY   CHAR(3)
ACCOUNTID  NUMBER

LIMITS
CURRENCY   CHAR(3)
MAX_AMOUNT NUMBER

EXCHANGE_RATES
AS_OF      DATE
CURRENCY   CHAR(3)
CURRENCY_VALUE NUMBER

```

```

Loop on ACCOUNTID values from one AREAID
Search TRANSACTIONS (last 30 days)
  Get MAX_AMOUNT for current CURRENCY
  If MONEY_AMOUNT >= MAX_AMOUNT
    Convert MONEY_AMOUNT
    Insert transaction id and converted amount
    into LOG_TABLE

```

This is pseudo-code for the business process.

```


Function getLimit(this_currency)
{
  select max_amount
  from limits
  where currency = this_currency;
}

Function convertAmount(this_amount,
                      this_currency,
                      the_date)
{
  select this_amount * currency_value
  from exchange_rates
  where as_of = the_date
  and currency = this_currency;
}

```

Two main stored functions were involved.

minutes



The process was taking far too much time. The correct approach for a DBA is to trace it, look for waits and see where most time is spent.

SQL ID : lnup7kcbvt072

```
select txid,money,amount,currency from transactions
where accountid=:1 and txdate >= to_date(:2, 'DD-MON-YYYY') - 30
order by txdate
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	270	0.01	0.02	0	0	0	0
Fetch	3252	38.13	39.67	3612	2317950	0	31029
total	3523	38.14	39.67	3612	2317950	0	31029

Rows Row Source Operation

```
134 SORT ORDER BY (cr=58 pr=3612 pw=3612 time=3 us cost=2497
size=2875 card=115)
134 TABLE ACCESS FULL TRANSACTIONS (cr=8585 pr=3612 pw=3612
time=80474 us cost=2496 size=2875 card=115)
```

missing index!

The trace file shows that we are scanning the 2,000,000 row table, which for most DBAs has an obvious answer.

SQL ID : gx2cn564cdsds

```
select max_amount from limits
where currency=:1
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	31029	0.74	0.77	0	0	0	0
Execute	31029	1.61	1.47	0	0	0	0
Fetch	31029	1.16	1.12	1	61330	0	30301
total	93087	3.51	3.38	1	61330	0	0301

Rows Row Source Operation


```
1 TABLE ACCESS BY INDEX ROWID LIMITS (cr=2 pr=1 pw=1 time=0 us
cost=1 size=9 card=1)
1 INDEX UNIQUE SCAN SYS_C009718 (cr=1 pr=0 pw=0 time=0 us cost=0
size=0 card=1) (object id 70976)
```

The other significant query represents only 10% of the first one (elapsed time smaller than CPU time is because of rounding errors)

```
CREATE INDEX MY_INDEX ON TRANSACTIONS (ACCOUNTID)
```

So let's create the missing index and ready ourselves for an impressive speed improvement.

minutes




There is some noticeable improvement, but not the "Wow!" kind of improvement that we were expecting. It's still far too slow. But there were TWO columns in the condition. Of course, we should have a composite index.

```
CREATE INDEX MY_INDEX ON TRANSACTIONS (ACCOUNTID,
                                         TXDATE)
```

There we go.


minutes



Same thing ...

Disappointment and head-scratching.

Trace file analysis (128 seconds)



In fact, when we compare measured times to times reported in the trace file, we notice that a lot of time is unaccounted for.

```
SQL ID : gx2cn564cdsds
select max_amount from limits
where currency=:1
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	31029	0.74	0.77	0	0	0	0
Execute	31029	1.61	1.47	0	0	0	0
Fetch	31029	1.16	1.12	1	61330	0	30301
total	93087	3.51	3.38	1	61330	0	0301

Rows Row Source Operation

```
1 TABLE ACCESS BY INDEX ROWID LIMITS (cr=2 pr=1 pw=1 time=0 us
cost=1 size=9 card=1)
1 INDEX UNIQUE SCAN SYS_C009718 (cr=1 pr=0 pw=0 time=0 us cost=0
size=0 card=1)(object id 70976)
```


What should have alerted us is this. There may be 20 currencies at most. Do we need to query limits THAT many times?

```

Function getLimit(this_currency)
{
  if (this_currency is new) {
    select max_amount
    from limits
    where currency = this_currency;
    store into cache
  }
  return value from cache
}

```

I have warned you against "look-up functions". This is one. Let's use a local cache and only run the query when we don't know the limit.



join ?

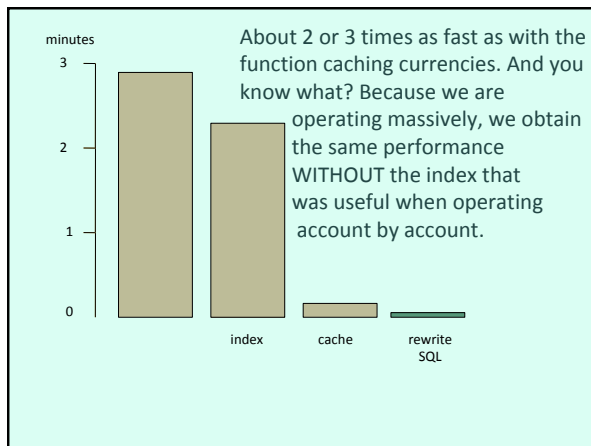
But rather than trying to improve the process by by-passing unnecessary queries, what about trying to rethink it globally? In fact, it doesn't need to be procedural. It can be a single SQL statement.

```

insert into log_table(txid,
                     converted_amount)
select x.txid,x.money_amount*e.currency_value
from (select t.txid,
            t.money_amount,
            t.currency
      from transactions t
      inner join accounts a
        on a.accountid = t.accountid
      inner join limits l
        on l.currency = t.currency
      where a.areaid = :the_area
            and t.txdate >= to_date(:the_date, 'DD-MON-YYYY') - 30
            and t.money_amount >= l.max_amount) x
      inner join exchange_rates e
        on e.currency = x.currency
where e.as_of = :the_date;

```

No look-up function, no loops, plain SQL (not even competition-grade SQL).



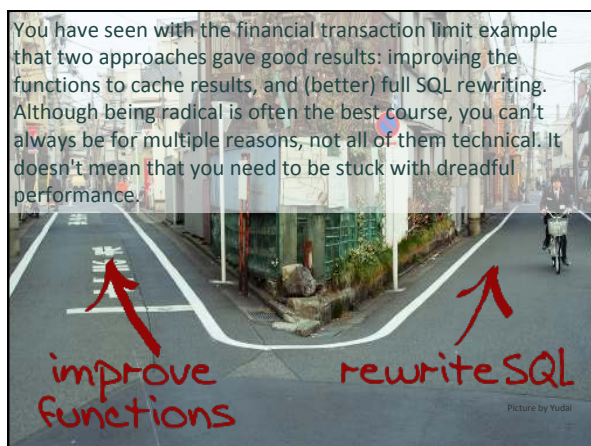
Lessons

Improving performance is not only adding indexes


Waits are useful for tuning, less so for refactoring

You can choose your approach

DON'T GET INTO THE WAY OF THE OPTIMIZER



I would also like to underline that it's not because one process revealed a problem that this process is the full reason for this problem. It may be just the straw that broke the camel (or donkey)'s back.



You should take a look at everything, even processes that had kept below the radar so far.

THE PROBLEM

The problem of many, many database applications can be explained with an analogy.

When they need to shop:

Developers take their car

Drive to the shop

Look for a parking place

Walk the aisles

Queue for paying

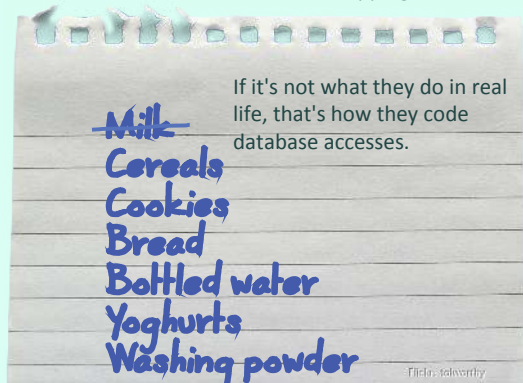
Look for their car

Drive back home


Store what they have bought



Then look for the next item on the shopping list.



Interestingly, improving EVERY point where you waited may not be the best solution.



Flickr: Taylorandayumi

**A good DBA approach
to performance issues**


isn't

**a good developer approach
to performance issues.**

You shouldn't be short-sighted in your approach and try to understand what the DBMS is getting and what it can do.


Seen by the DBMS

The DBMS only sees individual queries. It has no idea whatsoever about "flows".



Thousands of active users can be executing queries from the same program simultaneously.

Optimizer



Critically, the scope for the optimizer is ONE statement (and we have seen that optimizing, and getting stable performance, for one statement is hard enough)

```

for rec in (select deptno
            from emp
            where job = 'MANAGER')
loop
    select dname
    into v_deptname
    from dept
    where deptno = rec.deptno;
end loop

```

If you write your queries like this, the optimizer sees two independent queries that cannot be much improved. In effect, you are shoving a nested loop down the throat of the DBMS when a hash join might have been the smart way of retrieving data.

UNTUNABLE queries

This is something that you see very, very often, and that Object Relational Mappers seem to be particularly apt at generating: tons of queries where the only search criterion is the primary key. The problem is, if that primary key is system-generated as it often is, how did you find its value in the first place? You are obviously running (at least) two queries when one could have done the job.

Successive statements

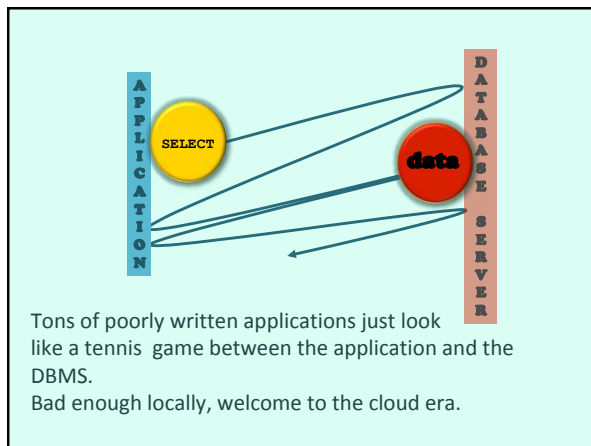
Another common mistake is when people execute in succession many statements that could have been run as one.

```

PROCEDURE p_hist_cust (p_txcode IN customers.txcode%TYPE,
                      p_seqnum IN OUT customers.seqnum%TYPE)
IS
BEGIN
    /* Insert into the customer history table */
    INSERT INTO histcust
    SELECT * FROM customers
    WHERE customers.txcode = p_txcode;
    /* Change the fields in the history table */
    UPDATE histcust SET histcust.indact = 'N'
    WHERE histcust.txcode = p_txcode
    AND histcust.seqnum = p_seqnum;
    /* Assign a new sequence number */
    p_seqnum := NVL(histcust.seqnum, 0) + 1;
EXCEPTION
    WHEN others THEN
        util.p_raise_error( ... );
END

```

You have here a real-life example. Why not inserting the right value directly instead of changing it afterwards? Can be done with a CASE ... END. Besides, select * is bad and trapping all errors dangerous.



Procedural logic



Far too few developers think in sets; and all the art of working with a database (and EVERYBODY is working with databases) is to switch from a procedural logic to a pure SQL logic. It's not because your procedural logic is embedded in a stored procedure that the sin is washed away.


```
for rec in (select deptno
            from emp
            where job = 'MANAGER')
loop
  select dname
  into v_deptname
  from dept
  where deptno = rec.deptno;
end loop
```

I have given this example earlier. It may be a bit extreme but when you look carefully enough you meet the pattern very often. People love cursors. Almost every database programming course focuses on them.

Nested loop
+ context switch / network latency

Is it the **BEST** we can do?

Loops are bad in a database environment. Very bad.



Why loops?

Sometimes there is a justification, or people think that there are justifications, for loops.

Input, output ?

OK.

Looping for writing to a file or sending data over a network, or in a procedural language that accesses the database is perfectly justified, because you are at the border between a world that knows sets and a world that (in the best of cases) only knows collections.

Transaction & Error Management ?

Question ...

Some people like to loop over changes to tables so as to be able to commit on a regular basis, and not have too big a transaction, or not have a big load that fails because of one wrong line of input. There may be some justification here, but don't take it for granted. A single transaction may be much faster, and restarting a process that has partly succeeded is always difficult (more than restarting from scratch).

`if ... then ...else`
in the loop?
Try harder

There are also cases when people loop to be able to implement conditional logic in the loop. Cases when it's required are very rare. Most often, conditions can be reported inside the SQL statement with CASE ... END constructs and UNION ALL queries.

Otherwise ...
No way!

These are about the only cases when the use of loops can (sometimes) be justified.

計利以聽，乃爲之勢，以佐其外。

Sun Tzu (6th century BC)
The Art of War, I,16

While heeding the profit of my counsel, avail yourself also of any helpful circumstances over and beyond the ordinary rules.